



Parking Lot Analysis

Distributed Systems WS2023/24

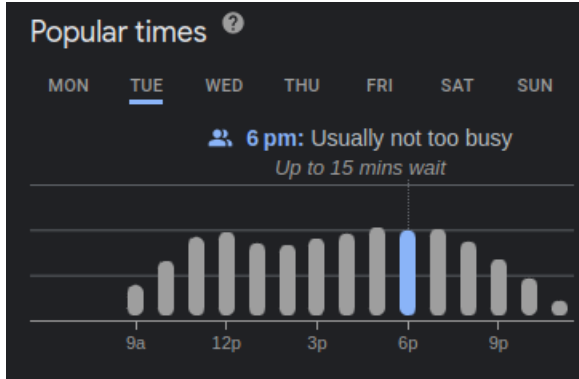
Simon Draxl, Luca Rahm, Evin Aydin, Alexander Montag

Content

- Motivation
- Workflow
- Live Demo
- Evaluation

Motivation: Popular Times for Restaurants

Googling a restaurant gives the following graph:



Source: [1]

Measure Busy Times for Parking Lots

Something similar for parking lots?


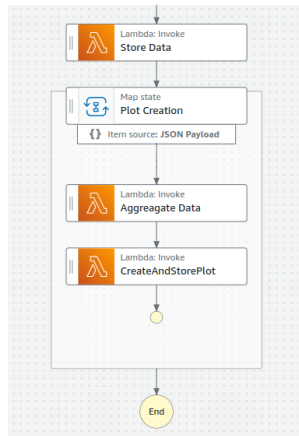
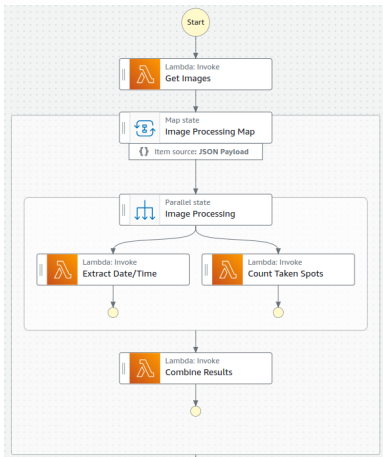


Image from the Dataset, source:
[2]

Desired output

Workflow



Workflow in Detail I

- **Get Images:** Retrieves the images from an S3 Bucket
- **Extract Data/Time:** Extracts date, time and day of week from an image name
- **Count Taken Spots:** Counts the taken spots for an image using Amazon Rekognition
- **Combine Results:** Combines the JSON output from the previous Lambda Functions
- **Store Data:** Stores the following data in a DynamoDB table (one table entry per timestamp):
 - Time
 - Date
 - Weekday
 - Taken spots

Workflow in Detail II

- Total spots
- **Aggregate Data:** Retrieves the data previously stores in the DynamoDB table and aggregates the data for each timestamp per weekday
- **Create and Store Plot:** Creates a plot for the average taken spots over the time of day for every weekday

Live Demo

Execution Times

For 120 images the execution time of ~ 1 min breaks down into the following parts:

- Get images: < 2 s
- Processing images: ~ 30 s
- Storing data in DynamoDB table: ~ 17 s
- Retrieving and aggregating data from table: < 1 s
- Creating and storing plot: ~ 11 s

Using concurrency limit of 5 parallel iterations

Execution Cost

One execution with 120 images leads to the following costs:

- Step functions: 0\$ (500 state transitions, 4000 per month are free)
- Lambda functions: 0\$ (only at a few 100,000 requests per month the lambda functions would exceed the free tier limit)
- Rekognition: 0.12\$

Daily executions for images every 5 min (96 images per day) leads to monthly execution costs of about 3.10\$

Does not include storage costs for DynamoDB and S3

Limitations of the Application



- Total number of parking spots needs to be an input and cannot be counted using Amazon Rekognition
- Key-value store does not allow easy aggregation → using SQLServer would be more efficient, but no access due to the student account
- Amazon Rekognition not entirely perfect when labeling cars, especially for images with a lot of cars (cars become very small in the images)



Thank you for your attention

Simon Draxl, Luca Rahm, Evin Aydin, Alexander Montag

References I

-  Google restaurant popular times.
<https://www.google.com/search?q=uni+cafe+innsbruck>.
Accessed: 22.01.2023.
-  Parking lot dataset.
<https://www.kaggle.com/datasets/ammarnassanalhajali/pklot-dataset>.
Accessed: 22.01.2023.